

Local and Distributed Matrices

Lesson Objectives

- After completing this lesson, you should be able to:
 - Understand local and distributed matrices
 - Create dense and sparse matrices
 - Create different types of distributed matrices

Local Matrices

- Natural extension of Vectors
- Row and column indices are 0-based integers and values are doubles
- Local matrices are stored on a single machine
- MLlib's matrices can be either dense or sparse
- Matrices are filled in column major order

Dense Matrices

- A "reshaped" dense Vector
- First two arguments specify dimensions of the matrix
- Entries are stored in a single double array

A Simple Dense Matrix

```
from pyspark.mllib.linalg import Matrix, Matrices
```

```
Matrices.dense(3, 2, [1, 3, 5, 2, 4, 6])
```

```
DenseMatrix(3, 2, [1.0, 3.0, 5.0, 2.0, 4.0, 6.0], False)
```

Sparse Matrices in Spark: Compressed Sparse Column (CSC) format

	0	1	2	3
0				
1		34		
2				
3			55	
4				

Rows: 5

Columns: 4

Column pointers: (0, 0, 1, 2, 2)

Row indices: (1, 3)

Non-zero values: (34, 55)

Sparse Matrix Example

```
Matrices.sparse(5, 4, [0,0,1,2,2], [1,2], [34,55])
```

```
SparseMatrix(5, 4, [0, 0, 1, 2, 2], [1, 2], [34.0, 55.0], False)
```

Distributed Matrices

- Distributed matrices are where Spark starts to deliver significant value
- They are stored in one or more **RDDs**
- Three types have been implemented:
 - **RowMatrix**
 - **IndexedRowMatrix**
 - **CoordinateMatrix**
- Conversions may require an expensive global shuffle

RowMatrix

- The most basic type of distributed matrix
- It has no meaningful row indices, being only a collection of feature vectors
- Backed by an **RDD** of its rows, where each row is a local vector

RowMatrix

- Assumes the number of columns is small enough to be stored in a local vector
- Can be easily created from an instance of `RDD[Vector]`

A Simple RowMatrix

```
from pyspark.rdd import RDD
from pyspark.mllib.linalg.distributed import RowMatrix
```

```
rows = sc.parallelize([Vectors.dense(1.0,2.0),
                        Vectors.dense(4.0,5.0),
                        Vectors.dense(7.0,8.0)])
```

A Simple RowMatrix

```
mat = RowMatrix(rows)
```

```
mat.numRows()
```

3L

```
mat.numCols()
```

2L

IndexedRowMatrix

- Similar to a `RowMatrix`
- But it has meaningful row indices, which can be used for identifying rows and executing joins
- Backed by an `RDD` of indexed rows, where each row is a tuple containing an index (long-typed) and a local vector
- Easily created from an instance of `RDD[IndexedRow]`
- Can be converted to a `RowMatrix` by calling `toRowMatrix()`

A Simple IndexedRowMatrix

```
from pyspark.mllib.linalg import Vector, Vectors
from pyspark.mllib.linalg.distributed import IndexedRow, IndexedRowMatrix
```

```
idx_rows = sc.parallelize([IndexedRow(0, Vectors.dense(1.0, 2.0)),
                           IndexedRow(1, Vectors.dense(4.0, 5.0)),
                           IndexedRow(2, Vectors.dense(7.0, 8.0))])
```

```
idx_mat = IndexedRowMatrix(idx_rows)
```

```
idx_mat.rows.collect()
```

```
[IndexedRow(0, [1.0, 2.0]), IndexedRow(1, [4.0, 5.0]), IndexedRow(2, [7.0, 8.0])]
```

CoordinateMatrix

- Should be used only when both dimensions are huge and the matrix is very sparse
- Backed by an **RDD** of matrix entries, where each entry is a tuple `(i: Long, j: Long, value: Double)` where:
 - i is the row index
 - j is the column index
 - value is the entry value

CoordinateMatrix

- Can be easily created from an instance of `RDD[MatrixEntry]`
- Can be converted to an `IndexedRowMatrix` with sparse rows by calling `toIndexedRowMatrix()`

A Simple CoordinateMatrix

```
from pyspark.mllib.linalg.distributed import MatrixEntry, CoordinateMatrix
```

```
entries = sc.parallelize([MatrixEntry(0,0,9.0),  
                             MatrixEntry(1,1,8.0),  
                             MatrixEntry(2,1,6.0)])
```

```
coord_mat = CoordinateMatrix(entries)
```

```
coord_mat.toIndexedRowMatrix().rows.collect()
```

```
[IndexedRow(0, (2,[0],[9.0])),  
 IndexedRow(1, (2,[1],[8.0])),  
 IndexedRow(2, (2,[1],[6.0]))]
```

Lesson Summary

- Having completed this lesson, you should be able to:
 - Understand local and distributed matrices
 - Create dense and sparse matrices
 - Create different types of distributed matrices