

# Evaluation

# Lesson Objectives

- After completing this lesson, you should be able to:
  - Evaluate binary classification algorithms using area under the ROC curve
  - Evaluate multiclass classification algorithms using several metrics
  - Evaluate regression algorithms using several metrics
  - Evaluate logistic and linear regression algorithms using summaries

# Evaluators

- Computes metrics from predictions
- Available Evaluators
  - BinaryClassificationEvaluator
  - MultiClassClassificationEvaluator
  - RegressionEvaluator

# BinaryClassificationEvaluator

- Evaluator for binary classification
- Expects two input columns: **rawPrediction** and **label**
- Supported metric: areaUnderROC

# BinaryClassificationEvaluator

```
from pyspark.ml.evaluation import BinaryClassificationEvaluator

predictionsLogR = logrModel.transform(testData)

evaluator = BinaryClassificationEvaluator().setLabelCol("label") \
                                             .setRawPredictionCol("rawPrediction") \
                                             .setMetricName("areaUnderROC")

roc = evaluator.evaluate(predictionsLogR)
print roc
```

1.0

# MulticlassClassificationEvaluator

- valuator for multiclass classification
- Expects two input columns: prediction and label
- Supported metrics:
  - F1 (default)
  - Precision
  - Recall
  - weightedPrecision
  - weightedRecall

# MulticlassClassificationEvaluator

```
from pyspark.ml.evaluation import MulticlassClassificationEvaluator

evaluator = MulticlassClassificationEvaluator().setLabelCol("indexedLabel") \
                                                .setPredictionCol("prediction") \
                                                .setMetricName("accuracy")

accuracy = evaluator.evaluate(predictionsRFC)

print "Test Error = %s" % (1.0 - accuracy)
```

Test Error = 0.02777777777778

# RegressionEvaluator

- Evaluator for regression
- Expects two input columns: prediction and label
- Supported metrics:
  - **rmse**: root mean squared error (default)
  - **mse**: mean squared error
  - **r2**: R2, the coefficient of determination
  - **mae**: mean absolute error



# A Simple Regression Evaluator

```
from pyspark.ml.evaluation import RegressionEvaluator

evaluator = RegressionEvaluator().setLabelCol("label") \
                                .setPredictionCol("prediction") \
                                .setMetricName("rmse")

rmse = evaluator.evaluate(predictionsRFR)

print "Root Mean Squared Error (RMSE) = %s" % rmse
```

Root Mean Squared Error (RMSE) = 0.057735026919

# LogisticRegressionSummary

- **LogisticRegressionSummary** accessible through summary attribute of a **LogisticRegressionModel**
- Summarizes the model over the training set

# BinaryLogisticRegressionSummary

- Supported metrics:
  - areaUnderROC**: area under the receiver operating characteristic (ROC) curve
  - fMeasureByThreshold**: dataframe with two fields (threshold, F-Measure) curve with  $\beta = 1$
  - pr**: precision-recall curve, dataframe containing two fields recall, precision with (0.0, 1.0) prepended to it
  - precisionByThreshold**: dataframe with two fields (threshold, precision) curve
  - recallByThreshold**: dataframe with two fields (threshold, recall) curve
  - roc**: receiver operating characteristic (ROC) curve, dataframe having two fields (FPR, TPR) with (0.0, 0.0) prepended and (1.0, 1.0) appended to it

# A Simple Logistic Regression Summary (1)

```
trainingSummaryLR = logitModel.summary  
trainingSummaryLR.areaUnderROC
```

1.0

```
fMeasure = trainingSummaryLR.fMeasureByThreshold  
fMeasure.show(3)
```

```
+-----+-----+  
|          threshold|          F-Measure|  
+-----+-----+  
|0.7948616368231948|0.05405405405405406|  
|0.7946216047546033|0.10526315789473684|  
|0.7944791865060398|0.15384615384615385|  
+-----+-----+  
only showing top 3 rows
```

# A Simple Logistic Regression Summary (2)

```
from pyspark.sql import functions as F

maxFMeasure = fMeasure.agg({"F-Measure": "max"}).head()[0]
print maxFMeasure
maxFMeasure = fMeasure.agg(F.max(F.col("F-Measure"))).head()[0]
print maxFMeasure
|
bestThreshold = fMeasure.where(F.col("F-Measure") == maxFMeasure).select("threshold").head()[0]
print bestThreshold
```

1.0

1.0

0.741090198576

# A Simple Logistic Regression Summary (3)

```
trainingSummaryLR.pr.show(3)
trainingSummaryLR.precisionByThreshold.show(3)
```

	recall	precision
	0.0	1.0
	0.02777777777777776	1.0
	0.05555555555555555	1.0

only showing top 3 rows

	threshold	precision
	0.7948616368231948	1.0
	0.7946216047546033	1.0
	0.7944791865060398	1.0

only showing top 3 rows

# A Simple Logistic Regression Summary (4)

```
trainingSummaryLR.recallByThreshold.show(3)
trainingSummaryLR.roc.show(3)
```

```
+-----+-----+
|          threshold|          recall|
+-----+-----+
|0.7948616368231948|0.02777777777777776|
|0.7946216047546033| 0.05555555555555555|
|0.7944791865060398| 0.08333333333333333|
+-----+-----+
only showing top 3 rows
```

```
+---+-----+
|FPR|          TPR|
+---+-----+
|0.0|          0.0|
|0.0|0.02777777777777776|
|0.0| 0.05555555555555555|
+---+-----+
only showing top 3 rows
```

# LinearRegressionSummary

- Accessible through `summary` attribute of a `LinearRegressionModel`
- Summarizes the model over the training set
- Supported metrics:
  - `explainedVariance`: explained variance regression score
  - `meanAbsoluteError`: mean absolute error (L1-norm loss)
  - `meanSquaredError`: mean squared error (quadratic loss)
  - `r2`:  $R^2$ , the coefficient of determination
  - `residuals`: residuals (label - predicted value)
  - `rootMeanSquaredError`: root mean squared error



# A Simple Linear Regression Summary (1)

```
trainingSummaryLLS = lrModel.summary  
  
print trainingSummaryLLS.explainedVariance  
  
print trainingSummaryLLS.meanAbsoluteError  
  
print trainingSummaryLLS.meanSquaredError  
  
print trainingSummaryLLS.r2
```

```
0.0725699709659  
0.227189149065  
0.0536653700314  
0.781931194793
```

# A Simple Linear Regression Summary (2)

```
trainingSummaryLLS.residuals.show(3)  
  
print trainingSummaryLLS.rootMeanSquaredError
```

```
+-----+  
|           residuals|  
+-----+  
|-0.23933577912497733|  
| -0.2785526569681104|  
|-0.24240274742471352|  
+-----+  
only showing top 3 rows  
  
0.231657872802
```

# Lesson Summary

- Having completed this lesson, you should be able to:
  - Evaluate binary classification algorithms using area under the ROC curve
  - Evaluate multiclass classification algorithms using several metrics
  - Evaluate regression algorithms using several metrics
  - Evaluate logistic and linear regression algorithms using summaries