

Linear Methods

Lesson Objectives

- After completing this lesson, you should be able to:
 - Understand the Pipelines API for Logistic Regression and Linear Least Squares
 - Perform classification with Logistic Regression
 - Perform regression with Linear Least Squares
 - Use regularization with Logistic Regression and Linear Least Squares

Linear Methods

- Logistic Regression
- Linear Least Squares

Logistic Regression

- Widely used to predict binary responses (classification)
- Can be generalized into multinomial logistic regression (multiclass)
 - for K possible outcomes, choose one outcome as "pivot"
 - the other $K-1$ outcomes can be separately regressed against the pivot outcome

Logistic regression advantages

- Has no tuning parameters
- Its prediction equation is simple and easy to implement

MLlib's implementation

- MLlib chooses first class (0) as the "pivot" class
- For multiclass classification problem, outputs a multinomial logistic regression model, containing K-1 binary logistic regression models regressed against the "pivot" class
- For a new data point, K-1 models are run and the class with largest probability is chosen as the predicted class
- Supported algorithms:
 - mini-batch gradient descent
 - L-BFGS (recommended for faster convergence)

Regularization

- L2 regularization: Ridge Regression (penalizes beta parameters by the square of their magnitude)
- L1 regularization: Lasso (penalizes beta parameters by their absolute value)
- Elastic net regularization combines L1 and L2, with a weight for each
 - equivalent to ridge regression (L2) if **alfa** set to 0
 - equivalent to Lasso (L1) if **alfa** set to 1

Elastic net regularization: parameters

- `elasticNetParam` corresponds to `alfa`
- `regParam` corresponds to `lambda`

	regularizer $R(\mathbf{w})$	gradient or sub-gradient
zero (unregularized)	0	0
L2	$\frac{1}{2} \ \mathbf{w}\ _2^2$	\mathbf{w}
L1	$\ \mathbf{w}\ _1$	$\text{sign}(\mathbf{w})$
elastic net	$\alpha \ \mathbf{w}\ _1 + (1 - \alpha) \frac{1}{2} \ \mathbf{w}\ _2^2$	$\alpha \text{sign}(\mathbf{w}) + (1 - \alpha) \mathbf{w}$

Example of Logistic Regression (1)

```
from pyspark.mllib.classification import LogisticRegressionWithLBFGS, LogisticRegressionModel
from pyspark.mllib.regression import LabeledPoint
```

```
!wget https://raw.githubusercontent.com/apache/spark/master/data/mllib/sample_svm_data.txt
```

```
--2016-09-26 08:55:31-- https://raw.githubusercontent.com/apache/spark/master/data/mllib/sample_svm_data.txt
```

```
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 151.101.12.133
```

```
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|151.101.12.133|:443... connected.
```

```
HTTP request sent, awaiting response... 200 OK
```

```
Length: 39474 (39K) [text/plain]
```

```
Saving to: 'sample_svm_data.txt.1'
```

```
100%[=====>] 39.474      --.-K/s   in 0,06s
```

```
2016-09-26 08:55:31 (670 KB/s) - 'sample_svm_data.txt.1' saved [39474/39474]
```

Example of Logistic Regression (2)

```
def parsePoint(line):  
    values = [float(x) for x in line.split(' ')]  
    return LabeledPoint(values[0], values[1:])
```

```
data = sc.textFile("sample_svm_data.txt")
```

```
parsedData = data.map(parsePoint)
```

```
parsedData.take(1)
```

```
[LabeledPoint(1.0, [0.0, 2.52078447202, 0.0, 0.0, 0.0, 2.00468443649, 2.00034729927, 0.0, 2.22838704274, 2.22838704274, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0])]
```

Example of Logistic Regression (3)

```
model = LogisticRegressionWithLBFGS.train(parsedData)

labelsAndPreds = parsedData.map(lambda p: (p.label, model.predict(p.features)))

trainErr = labelsAndPreds.filter(lambda (v, p): v != p).count() / float(parsedData.count())
print("Training Error = " + str(trainErr))
```

Training Error = 0.366459627329

Linear Least Squares

- Most common formulation for regression problems
- As in logistic regression, different types of regularization are possible:
 - no regularization: Ordinary Least Squares
 - L2 regularization: Ridge Regression
 - L1 regularization: Lasso
 - Elastic net
- Average loss = Mean Squared Error

Example of Linear Regression (1)

```
from pyspark.mllib.regression import LinearRegressionWithSGD, LinearRegressionModel
```

```
!wget https://raw.githubusercontent.com/apache/spark/master/data/mllib/ridge-data/lpsa.data
```

```
--2016-09-26 08:58:53-- https://raw.githubusercontent.com/apache/spark/master/data/mllib/ridge-data/lpsa.data
```

```
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 151.101.12.133
```

```
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|151.101.12.133|:443... connected.
```

```
HTTP request sent, awaiting response... 200 OK
```

```
Length: 10395 (10K) [text/plain]
```

```
Saving to: 'lpsa.data'
```

```
100%[=====>] 10.395      --.-K/s   in 0,001s
```

```
2016-09-26 08:58:54 (11,1 MB/s) - 'lpsa.data' saved [10395/10395]
```

Example of Linear Regression (2)

```
def parsePoint(line):  
    values = [float(x) for x in line.replace(',', ' ').split(' ')]  
    return LabeledPoint(values[0], values[1:])
```

```
data = sc.textFile("lpsa.data")  
parsedData = data.map(parsePoint)
```

```
parsedData.take(1)
```

```
[LabeledPoint(-0.4307829, [-1.63735562648, -2.00621178481, -1.86242597251, -1.02470580167, -0.522940888712, -0.863171185426, -1.04215728919, -0.864466507337])]
```

Example of Linear Regression (3)

```
model = LinearRegressionWithSGD.train(parsedData, iterations=100, step=0.00000001)

valuesAndPreds = parsedData.map(lambda p: (p.label, model.predict(p.features)))

MSE = valuesAndPreds.map(lambda (v, p): (v - p)**2).reduce(lambda x, y: x + y) / valuesAndPreds.count()

print("Mean Squared Error = " + str(MSE))
```

Mean Squared Error = 7.4510328101

Lesson Summary

- Having completed this lesson, you should be able to:

- Understand the Pipelines API for Logistic Regression and Linear Least Squares
- Perform classification with Logistic Regression
- Perform regression with Linear Least Squares
- Use regularization with Logistic Regression and Linear Least Squares